



## Introduction

This notebook analyzes customer behavior for that we have BigMarts Sales data collected in 2013 which is bifurcated in train (8523 records & 12 attributes) and test (5681 records & 11 attributes) data set, train data set has both independent and dependent variable(s) given below

- Item\_Identifier: Product ID
- Item\_Weight: Weight of Product
- Item\_Fat\_Content: Fat content of Product- Low/Regular
- Item\_Visibility: Parameter to know the visibility/reach of product
- Item\_Type: Category of Product
- Item\_MRP: Maximum Retail Price of the Product
- Outlet\_Identifier: Store ID
- Outlet\_Establishment\_Year: The Year in which store is established
- Outlet\_Size: Area-wise distribution of Stores- Low/Medium/High
- Outlet\_Location\_Type: Type of city in which outlet is located
- Outlet\_Type: Type of outlet - Grocery store or supermarket
- Item\_Outlet\_Sales: Sale price of product - The dependent variable to be predicted

## The Hypothesis

1. Locality: Outlet in populated locality should generate more revenue
2. Spending Capacity: Tier 1 should have more spending capacity than tier 2 and tier 3

3. Opening capacity: Item which have more opening capacity than other item have

3. Product Selection: Tier 1 should prefer low fat content food as they tend to be more aware of their health
4. Item Visibility: More visible Item should have more revenue generating power
5. Area: Stores which have early establishment could have higher outlet size
6. MRP: Consumers prefer reasonable product or Branded products

## Problem Statment

We need to analyse the dataset and come up with more insights and our main objective is to predict the Sales figure for the test dataset

## Approach

1. By applying Exploratory Data Analysis we will identify the relation between different attributes and evaluate meaningful information
2. By applying different supervised machine learning algorithms we will predict the sales for test dataset

## Aim of the Project

The main objective is to find the sales per product for each store and evaluate meaningful insights. Using this model, BigMart will try to understand different attributes of the product and apply them to increase their overall sales

## Part 1: Data Preprocessing

### Importing the Librabies and Dataset

```

In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker
# For example, here's several helpful packages to load

import os #paths to file
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import warnings# warning filter

#ploting libraries
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set_palette('husl')

#feature engineering
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder

#train test split
from sklearn.model_selection import train_test_split

#metrics
from sklearn.metrics import mean_absolute_error as MAE
from sklearn.metrics import mean_squared_error as MSE
from sklearn.metrics import r2_score as R2
from sklearn.model_selection import cross_val_score as CVS

#ML models
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import Lasso

#default theme and settings
sns.set(context='notebook', style='darkgrid', palette='deep', font='sans-serif',
pd.options.display.max_columns

#warning hadle
warnings.filterwarnings("always")
warnings.filterwarnings("ignore")

## Display all the columns of the dataframe
pd.pandas.set_option('display.max_columns',None)

## Display all the rows of the dataframe
#pd.pandas.set_option('display.max_rows',None)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list a

import os
for dirname, _, filenames in os.walk('/kaggle/input'):

```

```

for filename in filenames:
    print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets
# You can also write temporary files to /kaggle/temp/, but they won't be saved ou
/kaggle/input/bigmart-sales-data/train.csv
/kaggle/input/bigmart-sales-data/Test.csv

```

```

In [2]: ## Loading the train dataset
dataset = pd.read_csv('../input/bigmart-sales-data/Train.csv')

```

```

## print shape of dataset with rows and columns
print(dataset.shape)

```

```
(8523, 12)
```

## Data Exploration

```

In [3]: dataset.head()

```

```

Out[3]:

```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Iden
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OU
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OU
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OU
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OU
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OU

```

In [4]: # check the columns
dataset.columns

```

```

Out[4]: Index(['Item_Identifier', 'Item_Weight', 'Item_Fat_Content', 'Item_Visibility',
              'Item_Type', 'Item_MRP', 'Outlet_Identifier',
              'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
              'Outlet_Type', 'Item_Outlet_Sales'],
              dtype='object')

```

In [5]: *# check the information about the dataset*  
dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       8523 non-null   object
1   Item_Weight                          7060 non-null   float64
2   Item_Fat_Content                     8523 non-null   object
3   Item_Visibility                      8523 non-null   float64
4   Item_Type                           8523 non-null   object
5   Item_MRP                            8523 non-null   float64
6   Outlet_Identifier                    8523 non-null   object
7   Outlet_Establishment_Year           8523 non-null   int64
8   Outlet_Size                         6113 non-null   object
9   Outlet_Location_Type                8523 non-null   object
10  Outlet_Type                          8523 non-null   object
11  Item_Outlet_Sales                   8523 non-null   float64
dtypes: float64(4), int64(1), object(7)
memory usage: 799.2+ KB
```

In [6]: *# Check the name of columns which contain string*  
dataset.select\_dtypes(include='object').columns

Out[6]: Index(['Item\_Identifier', 'Item\_Fat\_Content', 'Item\_Type', 'Outlet\_Identifier',  
'Outlet\_Size', 'Outlet\_Location\_Type', 'Outlet\_Type'],  
dtype='object')

In [7]: *# Check the no. of columns which contain string*  
len(dataset.select\_dtypes(include='object').columns)

Out[7]: 7

In [8]: *# Check the name of columns which contain numerical value*  
dataset.select\_dtypes(include=['int64', 'float64']).columns

Out[8]: Index(['Item\_Weight', 'Item\_Visibility', 'Item\_MRP',  
'Outlet\_Establishment\_Year', 'Item\_Outlet\_Sales'],  
dtype='object')

In [9]: *# Check the no. of columns which contain numerical value*  
len(dataset.select\_dtypes(include=['int64', 'float64']).columns)

Out[9]: 5

```
In [10]: # statistical summary
dataset.describe()
```

```
Out[10]:
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	7060.000000	8523.000000	8523.000000	8523.000000	8523.000000
mean	12.857645	0.066132	140.992782	1997.831867	2181.288914
std	4.643456	0.051598	62.275067	8.371760	1706.499616
min	4.555000	0.000000	31.290000	1985.000000	33.290000
25%	8.773750	0.026989	93.826500	1987.000000	834.247400
50%	12.600000	0.053931	143.012800	1999.000000	1794.331000
75%	16.850000	0.094585	185.643700	2004.000000	3101.296400
max	21.350000	0.328391	266.888400	2009.000000	13086.964800

## Dealing with missing data

```
In [11]: dataset.isnull().values.any()
```

```
Out[11]: True
```

```
In [12]: dataset.isnull().values.sum()
```

```
Out[12]: 3873
```

```
In [13]: dataset.isnull().sum()
```

```
Out[13]: Item_Identifier      0
Item_Weight      1463
Item_Fat_Content      0
Item_Visibility      0
Item_Type      0
Item_MRP      0
Outlet_Identifier      0
Outlet_Establishment_Year      0
Outlet_Size      2410
Outlet_Location_Type      0
Outlet_Type      0
Item_Outlet_Sales      0
dtype: int64
```


```
In [14]: # columns with null values
dataset.columns[dataset.isnull().any()]
```

```
Out[14]: Index(['Item_Weight', 'Outlet_Size'], dtype='object')
```

```
In [15]: len(dataset.columns[dataset.isnull().any()])
```

```
Out[15]: 2
```

```
In [16]: # null values with heatmap  
plt.figure(figsize=(16,9))  
sns.heatmap(dataset.isnull())  
plt.show()
```



```
In [17]: null_percent = dataset.isnull().sum() / dataset.shape[0] * 100


# (missing values / total values) * 100

null_percent
```

```
Out[17]: Item_Identifier      0.000000
Item_Weight      17.165317
Item_Fat_Content  0.000000
Item_Visibility  0.000000
Item_Type        0.000000
Item_MRP         0.000000
Outlet_Identifier 0.000000
Outlet_Establishment_Year 0.000000
Outlet_Size      28.276428
Outlet_Location_Type 0.000000
Outlet_Type      0.000000
Item_Outlet_Sales 0.000000
dtype: float64
```

```
In [18]: plt.figure(figsize=(8,5))
sns.boxplot('Item_Weight',data=dataset)
```

```
Out[18]: <AxesSubplot:xlabel='Item_Weight'>
```



**Box Plot suggest we dont have any outlier and hence we can change missing values with 'Mean'**

```
In [19]: dataset['Item_Weight'] = dataset['Item_Weight'].fillna(dataset['Item_Weight'].mean())
```

Since the Outlet\_Size is a categorical variable we can change this missing values to "Mode"(Most Repeated Value)



```
In [20]: dataset['Outlet_Size'] = dataset['Outlet_Size'].fillna(dataset['Outlet_Size'].mode()[0])
```

```
In [21]: dataset.isnull().values.any()
```

```
Out[21]: False
```

```
In [22]: len(dataset.columns[dataset.isnull().any()])
```

```
Out[22]: 0
```

## Cleaning the Data

```
In [23]: dataset['Item_Identifier'].value_counts()
```

```
Out[23]: FDG33      10
         FDW13      10
         NCB18       9
         DRE49       9
         FDX20       9
         ..
         FDQ60       1
         FDT35       1
         FDC23       1
         FDE52       1
         DRF48       1
         Name: Item_Identifier, Length: 1559, dtype: int64
```

```
In [24]: dataset['Item_Fat_Content'].value_counts()
```

```
Out[24]: Low Fat      5089
         Regular      2889
         LF           316
         reg          117
         low fat      112
         Name: Item_Fat_Content, dtype: int64
```

Some of 'Low Fat' values mis-coded as 'low fat' and 'LF'. Also, some of 'Regular' are mentioned as 'regular'. We need to fix them

```
In [25]: dataset['Item_Fat_Content'].replace(['low fat', 'LF', 'reg'], ['Low Fat', 'Low Fat', 'Regular'])
```

```
In [26]: dataset['Item_Fat_Content'].value_counts()
```

```
Out[26]: Low Fat      5517
         Regular      3006
         Name: Item_Fat_Content, dtype: int64
```

```
In [27]: dataset['Item_Type'].value_counts()
```

```
Out[27]: Fruits and Vegetables    1232
         Snack Foods              1200
         Household                910
         Frozen Foods             856
         Dairy                   682
         Canned                  649
         Baking Goods            648
         Health and Hygiene      520
         Soft Drinks             445
         Meat                    425
         Breads                  251
         Hard Drinks             214
         Others                  169
         Starchy Foods          148
         Breakfast              110
         Seafood                 64
         Name: Item_Type, dtype: int64
```

```
In [28]: dataset['Outlet_Identifier'].value_counts()
```

```
Out[28]: OUT027    935
         OUT013    932
         OUT049    930
         OUT046    930
         OUT035    930
         OUT045    929
         OUT018    928
         OUT017    926
         OUT010    555
         OUT019    528
         Name: Outlet_Identifier, dtype: int64
```

```
In [29]: dataset['Outlet_Size'].value_counts()
```

```
Out[29]: Medium    5203
         Small     2388
         High       932
         Name: Outlet_Size, dtype: int64
```

```
In [30]: dataset['Outlet_Location_Type'].value_counts()
```

```
Out[30]: Tier 3    3350
         Tier 2    2785
         Tier 1    2388
         Name: Outlet_Location_Type, dtype: int64
```

```
In [31]: dataset['Outlet_Type'].value_counts()
```

```
Out[31]: Supermarket Type1    5577
Grocery Store    1083
Supermarket Type3    935
Supermarket Type2    928
Name: Outlet_Type, dtype: int64
```

*We will convert "Outlet\_Establishment\_Year" to Age of the Store to get more meaning from the data*

```
In [32]: dataset['Years_Established'] = dataset['Outlet_Establishment_Year'].apply(lambda
dataset = dataset.drop(columns=['Outlet_Establishment_Year'])
dataset.head()
```

```
Out[32]:
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Iden
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OU
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OU
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OU
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OU
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OU

## Part 2: Exploratory Data Analysis

### A] Univariate Analysis


```
In [33]: # Check the name of coloumns which contain string
dataset.select_dtypes(include='object').columns
```

```
Out[33]: Index(['Item_Identifier', 'Item_Fat_Content', 'Item_Type', 'Outlet_Identifier',
'Outlet_Size', 'Outlet_Location_Type', 'Outlet_Type'],
dtype='object')
```

#### 1) Item Fat Content

```
In [34]: plt.figure(figsize=(8,5))  
sns.countplot('Item_Fat_Content',data=dataset,palette='Set1')
```

```
Out[34]: <AxesSubplot:xlabel='Item_Fat_Content', ylabel='count'>
```




Observation: People bought more Low Fat Items

## 2) Item Type

```
In [35]: plt.figure(figsize=(24,8))  
sns.countplot('Item_Type',data=dataset,palette='twilight')
```

```
Out[35]: <AxesSubplot:xlabel='Item_Type', ylabel='count'>
```




Observation: People bought more Fruits and Vegetables

### 3) Outlet Size

```
In [36]: plt.figure(figsize=(8,5))  
sns.countplot('Outlet_Size',data=dataset,palette='Set2')
```

```
Out[36]: <AxesSubplot:xlabel='Outlet_Size', ylabel='count'>
```

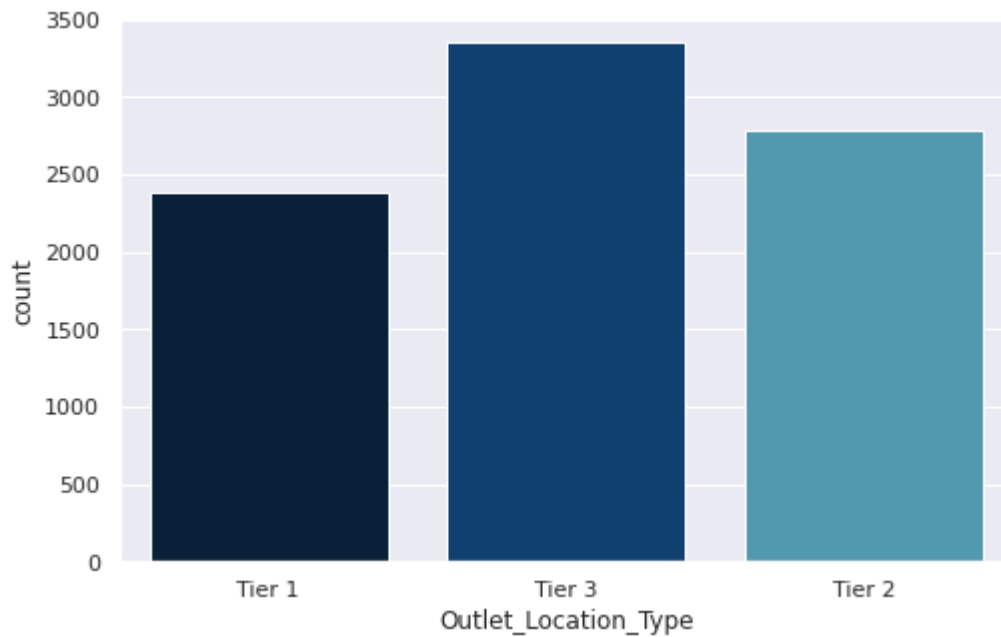


Observation: We have more Medium Outlets

### 4) Outlet Location

```
In [37]: plt.figure(figsize=(8,5))  
sns.countplot('Outlet_Location_Type',data=dataset,palette='ocean')
```

```
Out[37]: <AxesSubplot:xlabel='Outlet_Location_Type', ylabel='count'>
```




Observation: Maximum outlets in Tier 3 cities

## 5) Outlet Type

```
In [38]: plt.figure(figsize=(8,5))  
sns.countplot('Outlet_Type',data=dataset,palette='spring')
```

```
Out[38]: <AxesSubplot:xlabel='Outlet_Type', ylabel='count'>
```



Observation: Maximum supermarket are of Type 1

## B] Bivariate Analysis


```
In [39]: # Check the name of coloumns which contain string  
dataset.select_dtypes(include='object').columns
```

```
Out[39]: Index(['Item_Identifier', 'Item_Fat_Content', 'Item_Type', 'Outlet_Identifier',  
               'Outlet_Size', 'Outlet_Location_Type', 'Outlet_Type'],  
              dtype='object')
```

### 1) Item Fat Content to Item Outlet Sales

```
In [40]: plt.figure(figsize=(8,5))  
sns.barplot('Item_Fat_Content','Item_Outlet_Sales',data=dataset,palette='summer')
```

```
Out[40]: <AxesSubplot:xlabel='Item_Fat_Content', ylabel='Item_Outlet_Sales'>
```




Observation: Low Fat and Regular both are contributing equally to the revenue generation

## 2) Item Visibility to Item Outlet Sales



```
In [41]: plt.figure(figsize=(8,5))  
plt.scatter('Item_Visibility','Item_Outlet_Sales',data=dataset)
```

```
Out[41]: <matplotlib.collections.PathCollection at 0x7f9ef6f30290>
```




Observation: Here we have interesting observation, where the visibility of Items is Zero, which suggest those items kept behind in shelf and almost have no visibility can also be sold. This shows Consumer tend to search for their own products

### 3) Item Type to Item Outlet Sales

```
In [42]: plt.figure(figsize=(10,8))  
sns.barplot(y='Item_Type',x='Item_Outlet_Sales',data=dataset,palette='spring')
```

```
Out[42]: <AxesSubplot:xlabel='Item_Outlet_Sales', ylabel='Item_Type'>
```




Observation: Although Fruits and Vegetables unit sold are high, however revenue generated by Seafood is much higher, so we have to focus more on such products

#### 4) Item MRP to Item Outlet Sales

```
In [43]: plt.figure(figsize=(8,5))  
plt.scatter(y='Item_Outlet_Sales',x='Item_MRP',data=dataset)  
plt.xlabel('Item MRP')  
plt.ylabel('Item Outlet Sales')
```

```
Out[43]: Text(0, 0.5, 'Item Outlet Sales')
```




Observation: Items with higher MRP are sold maximum

## 5) Outlet Size to Item Outlet Sales

```
In [44]: plt.figure(figsize=(8,5))  
sns.barplot(x='Outlet_Size',y='Item_Outlet_Sales',data=dataset,palette='mako')
```

```
Out[44]: <AxesSubplot:xlabel='Outlet_Size', ylabel='Item_Outlet_Sales'>
```




Observation: Medium and High size outlet have maximum revenue generation power

## 6) Outlet Location to Item Outlet Sales

```
In [45]: plt.figure(figsize=(8,5))  
sns.barplot(x='Outlet_Location_Type',y='Item_Outlet_Sales',data=dataset,palette='mako')
```

```
Out[45]: <AxesSubplot:xlabel='Outlet_Location_Type', ylabel='Item_Outlet_Sales'>
```




Observation: Tier 2 & 3 have more revenue generation power although we have maximum number of outlet in Tier 3 cities so it justifies the number

## C] Multivariate Analysis

### 1) Item Type by Item Fat Content to Item Outlet Sales

```
In [46]: plt.figure(figsize=(25,5))
sns.barplot('Item_Type', 'Item_Outlet_Sales', hue='Item_Fat_Content', data=dataset,
plt.legend()
```

Out[46]: <matplotlib.legend.Legend at 0x7f9ef70749d0>

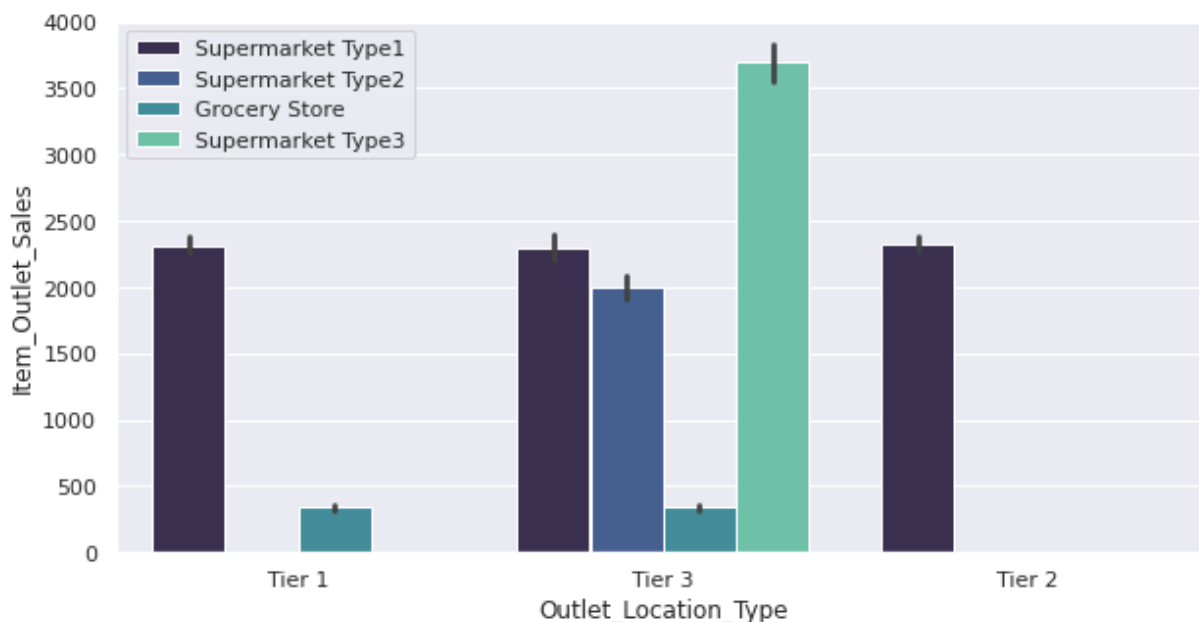


Observation: Mostly we have equal revenue generation from Fat and Regular Food Items

### 2) Outlet Location Type by Outlet Type to Item Outlet Sales

```
In [47]: plt.figure(figsize=(10,5))
sns.barplot('Outlet_Location_Type', 'Item_Outlet_Sales', hue='Outlet_Type', data=dataset,
plt.legend()
```

Out[47]: <matplotlib.legend.Legend at 0x7f9ef6dddad0>




Observation: Here we have interesting observation, where the visibility of Items is Zero, which suggest those items kept behind in shelf and almost have no visibility can also be sold. This show Consumer tend to search for their own products

## Distplot

In [48]: *# distplot of the target variable*

```
plt.figure(figsize=(16,9))
bar = sns.distplot(dataset['Item_Outlet_Sales'])
bar.legend(["Skewness: {:.2f}".format(dataset['Item_Outlet_Sales'].skew())])
plt.show()
```



## Correlation matrix


In [49]: `dataset_2 = dataset.drop(columns='Item_Outlet_Sales')`

In [50]: `dataset_2.shape`


Out[50]: (8523, 11)

```
In [51]: dataset_2.corrwith(dataset['Item_Outlet_Sales']).plot.bar(  
        figsize=(16,9), title='Correlated with SalePrice', grid=True  
        )
```

```
Out[51]: <AxesSubplot:title={'center':'Correlated with SalePrice'}>
```



```
In [52]: # heatmap
plt.figure(figsize=(9, 9))
ax = sns.heatmap(data=dataset.corr(), cmap='coolwarm', annot=True, linewidths=2)
```





## Part 3) Feature Engineering

### Label Encoding

```
In [53]: #feature engineering
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
var_mod = ['Outlet_Identifier', 'Item_Type']

for i in var_mod:
    dataset[i] = le.fit_transform(dataset[i])
```

```
In [54]: dataset.head()
```

```
Out[54]:
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Iden
0	FDA15	9.30	Low Fat	0.016047	4	249.8092	
1	DRC01	5.92	Regular	0.019278	14	48.2692	
2	FDN15	17.50	Low Fat	0.016760	10	141.6180	
3	FDX07	19.20	Regular	0.000000	6	182.0950	
4	NCD19	8.93	Low Fat	0.000000	9	53.8614	

### One Hot Encoding

```
In [55]: dataset = dataset.drop(columns=['Item_Identifier'])
dataset.head()
```

```
Out[55]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Siz
0	9.30	Low Fat	0.016047	4	249.8092	9	Mediu
1	5.92	Regular	0.019278	14	48.2692	3	Mediu
2	17.50	Low Fat	0.016760	10	141.6180	9	Mediu
3	19.20	Regular	0.000000	6	182.0950	0	Mediu
4	8.93	Low Fat	0.000000	9	53.8614	1	Hig

```
In [56]: #feature engineering
from sklearn.preprocessing import OneHotEncoder
```

```
In [57]: dataset = pd.get_dummies(data=dataset, drop_first=True)
dataset.shape
```

```
Out[57]: (8523, 15)
```

```
In [58]: dataset.head()
```

```
Out[58]:
```

	Item_Weight	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Item_Outlet_Sales	Years_Es
0	9.30	0.016047	4	249.8092	9	3735.1380	
1	5.92	0.019278	14	48.2692	3	443.4228	
2	17.50	0.016760	10	141.6180	9	2097.2700	
3	19.20	0.000000	6	182.0950	0	732.3800	
4	8.93	0.000000	9	53.8614	1	994.7052	

## Removing Skewness

Skewness in variables is undesirable for predictive modeling. Some machine learning methods assume normally distributed data and a skewed variable can be transformed by taking its log, square root, or cube root so as to make its distribution as close to normal distribution as possible. In our data, variables Item\_Visibility is highly skewed. So, we will treat skewness with the help of log transformation.

```
In [59]: #dataset['Item_Visibility'] = np.log(dataset['Item_Visibility'])
```

```
In [60]: #dataset.head()
```

*Skewness not taken into account as it was hindering the performance*

## Splitting the dataset

```
In [61]: # independ variables / matrix of features  
x = dataset.drop(columns='Item_Outlet_Sales')
```

```
In [62]: # target variable / dependent variable  
y = dataset['Item_Outlet_Sales']
```

```
In [63]: from sklearn.model_selection import train_test_split
```

```
In [64]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_s
```

```
In [65]: x_train.shape
```

```
Out[65]: (6818, 14)
```

```
In [66]: y_train.shape
```

```
Out[66]: (6818,)
```

```
In [67]: x_test.shape
```

```
Out[67]: (1705, 14)
```

```
In [68]: y_test.shape
```

```
Out[68]: (1705,)
```

## Feature scaling

```
In [69]: features= ['Item_Weight', 'Item_Fat_Content', 'Item_Visibility', 'Item_Type', 'Item_M
```



## Part 4: Building the model

```
In [70]: #metrics
from sklearn.metrics import mean_absolute_error as MAE
from sklearn.metrics import mean_squared_error as MSE
from sklearn.metrics import r2_score as R2
from sklearn.model_selection import cross_val_score as CVS
```

```
In [71]: #ML models
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import Lasso
from xgboost import XGBRFRegressor
from sklearn.linear_model import Ridge
```

```
In [72]: def cross_val(model_name,model,x,y,cv):

    scores = CVS(model, x, y, cv=cv)
    print(f'{model_name} Scores:')
    for i in scores:
        print(round(i,2))
    print(f'Average {model_name} score: {round(scores.mean(),4)}')
```

## 1) Multiple Linear Regressor

```
In [73]: #ML model
from sklearn.linear_model import LinearRegression

#model
regressor_mlr = LinearRegression()

#fit
regressor_mlr.fit(x_train, y_train)

#predict
y_pred = regressor_mlr.predict(x_test)

#score variables
LR_MAE = round(MAE(y_test, y_pred),2)
LR_MSE = round(MSE(y_test, y_pred),2)
LR_R_2 = round(R2(y_test, y_pred),4)
LR_CS = round(CVS(regressor_mlr, x, y, cv=5).mean(),4)

print(f" Mean Absolute Error: {LR_MAE}\n")
print(f" Mean Squared Error: {LR_MSE}\n")
print(f" R^2 Score: {LR_R_2}\n")
cross_val(regressor_mlr,LinearRegression(),x,y,5)
```

Mean Absolute Error: 851.52

Mean Squared Error: 1275654.43

R^2 Score: 0.5642

LinearRegression() Scores:

0.57

0.56

0.55

0.57

0.57

Average LinearRegression() score: 0.5613

```
In [74]: Linear_Regression=pd.DataFrame({'y_test':y_test,'prediction':y_pred})
Linear_Regression.to_csv("Linear Regression.csv")
```

## 2) Random Forest Regressor

```
In [75]: #ML model
from sklearn.ensemble import RandomForestRegressor

#model
regressor_rf = RandomForestRegressor(n_estimators=200,max_depth=5, min_samples_le

#fit
regressor_rf.fit(x_train, y_train)

#predict
y_pred = regressor_rf.predict(x_test)

#score variables
RFR_MAE = round(MAE(y_test, y_pred),2)
RFR_MSE = round(MSE(y_test, y_pred),2)
RFR_R_2 = round(R2(y_test, y_pred),4)
RFR_CS = round(CVS(regressor_rf, x, y, cv=5).mean(),4)

print(f" Mean Absolute Error: {RFR_MAE}\n")
print(f" Mean Squared Error: {RFR_MSE}\n")
print(f" R^2 Score: {RFR_R_2}\n")
cross_val(regressor_rf,RandomForestRegressor(),x,y,5)
```

Mean Absolute Error: 780.11

Mean Squared Error: 1200066.53

R^2 Score: 0.59

RandomForestRegressor(max\_depth=5, min\_samples\_leaf=100, n\_estimators=200,  
n\_jobs=4, random\_state=101) Scores:

0.57

0.53

0.53

0.55

0.57

Average RandomForestRegressor(max\_depth=5, min\_samples\_leaf=100, n\_estimators=200,  
n\_jobs=4, random\_state=101) score: 0.5504

```
In [76]: Random_Forest_Regressor=pd.DataFrame({'y_test':y_test,'prediction':y_pred})
Random_Forest_Regressor.to_csv("Random Forest Regressor.csv")
```

### 3) Lasso Regressor

```
In [77]: #ML model
from sklearn.linear_model import Lasso

#model
regressor_ls = Lasso(alpha = 0.05)
#fit
regressor_ls.fit(x_train,y_train)

#predict
y_pred = regressor_ls.predict(x_test)

#score variables
LS_MAE = round(MAE(y_test, y_pred),2)
LS_MSE = round(MSE(y_test, y_pred),2)
LS_R_2 = round(R2(y_test, y_pred),4)
LS_CS = round(CVS(regressor_ls, x, y, cv=5).mean(),4)

print(f" Mean Absolute Error: {LS_MAE}\n")
print(f" Mean Squared Error: {LS_MSE}\n")
print(f" R^2 Score: {LS_R_2}\n")
cross_val(regressor_ls,Lasso(alpha = 0.05),x,y,5)
```

Mean Absolute Error: 851.31

Mean Squared Error: 1275582.84

R^2 Score: 0.5642

Lasso(alpha=0.05) Scores:

0.57

0.56

0.55

0.57

0.57

Average Lasso(alpha=0.05) score: 0.5613

```
In [78]: Lasso_Regressor=pd.DataFrame({'y_test':y_test,'prediction':y_pred})
Lasso_Regressor.to_csv("Lasso Regressor.csv")
```

## 4) XGBoost Regressor

```

In [79]: #ML model
from xgboost import XGBRFRegressor

#model
regressor_xgb = XGBRFRegressor()

#fit
regressor_xgb.fit(x_train, y_train)

#predict
y_pred = regressor_xgb.predict(x_test)

#score variables
XGB_MAE = round(MAE(y_test, y_pred),2)
XGB_MSE = round(MSE(y_test, y_pred),2)
XGB_R_2 = round(R2(y_test, y_pred),4)
XGB_CS = round(CVS(regressor_xgb, x, y, cv=5).mean(),4)

print(f" Mean Absolute Error: {XGB_MAE}\n")
print(f" Mean Squared Error: {XGB_MSE}\n")
print(f" R^2 Score: {XGB_R_2}\n")
cross_val(regressor_xgb,XGBRFRegressor(alpha = 0.05),x,y,5)

```

Mean Absolute Error: 773.39

Mean Squared Error: 1194035.65

R^2 Score: 0.592

XGBRFRegressor(base\_score=0.5, booster='gbtree', colsample\_bylevel=1, colsample\_bytree=1, gamma=0, gpu\_id=-1, importance\_type='gain', interaction\_constraints='', max\_delta\_step=0, max\_depth=6, min\_child\_weight=1, missing=nan, monotone\_constraints='()', n\_estimators=100, n\_jobs=4, num\_parallel\_tree=100, objective='reg:squarederror', random\_state=0, reg\_alpha=0, scale\_pos\_weight=1, tree\_method='exact', validate\_parameters=1, verbosity=None) Scores:

0.61

0.58

0.57

0.6

0.61

Average XGBRFRegressor(base\_score=0.5, booster='gbtree', colsample\_bylevel=1, colsample\_bytree=1, gamma=0, gpu\_id=-1, importance\_type='gain', interaction\_constraints='', max\_delta\_step=0, max\_depth=6, min\_child\_weight=1, missing=nan, monotone\_constraints='()', n\_estimators=100, n\_jobs=4, num\_parallel\_tree=100, objective='reg:squarederror', random\_state=0, reg\_alpha=0, scale\_pos\_weight=1, tree\_method='exact', validate\_parameters=1, verbosity=None) score: 0.5935

```

In [80]: XGBoost_Regressor=pd.DataFrame({'y_test':y_test,'prediction':y_pred})
XGBoost_Regressor.to_csv("XGBoost_Regressor.csv")

```



## 5) Ridge Regressor

```
In [81]: #ML model
from sklearn.linear_model import Ridge

#model
regressor_rd = Ridge(normalize=True)
#fit
regressor_rd.fit(x_train,y_train)

#predict
y_pred = regressor_ls.predict(x_test)

#score variables
RD_MAE = round(MAE(y_test, y_pred),2)
RD_MSE = round(MSE(y_test, y_pred),2)
RD_R_2 = round(R2(y_test, y_pred),4)
RD_CS = round(CVS(regressor_rd, x, y, cv=5).mean(),4)

print(f" Mean Absolute Error: {RD_MAE}\n")
print(f" Mean Squared Error: {RD_MSE}\n")
print(f" R^2 Score: {RD_R_2}\n")
cross_val(regressor_rd,Ridge(normalize=True),x,y,5)
```

Mean Absolute Error: 851.31

Mean Squared Error: 1275582.84

R^2 Score: 0.5642

Ridge(normalize=True) Scores:

0.38

0.38

0.38

0.37

0.38

Average Ridge(normalize=True) score: 0.376

```
In [82]: Ridge_Regressor=pd.DataFrame({'y_test':y_test,'prediction':y_pred})
Ridge_Regressor.to_csv("Ridge Regressor.csv")
```

## Conclusion

```
In [83]: MAE= [LR_MAE,RFR_MAE,LS_MAE,XGB_MAE,RD_MAE]
MSE= [LR_MSE,RFR_MSE,LS_MSE,XGB_MSE,RD_MSE]
R_2= [LR_R_2,RFR_R_2,LS_R_2,XGB_R_2,RD_R_2]
Cross_score= [LR_CS,RFR_CS,LS_CS,XGB_CS,RD_CS]

Models = pd.DataFrame({
    'Models': ["Linear Regression", "Random Forest Regressor", "Lasso Regressor", ">
    'MAE': MAE, 'MSE': MSE, 'R^2':R_2, 'Cross Validation Score':Cross_score})
Models.sort_values(by='MAE', ascending=True)
```

```
Out[83]:
```

	Models	MAE	MSE	R^2	Cross Validation Score
3	XGBoost Regressor	773.39	1194035.65	0.5920	0.5935
1	Random Forest Regressor	780.11	1200066.53	0.5900	0.5948
2	Lasso Regressor	851.31	1275582.84	0.5642	0.5613
4	Ridge Regressor	851.31	1275582.84	0.5642	0.3760
0	Linear Regression	851.52	1275654.43	0.5642	0.5613

## Realizations

1. XGBoost Regressor and Random Forest are best performing Models, we can use both to check on test data set and find out which perform better
2. MRP has huge correlation with the Outlet Sales
3. For better performance we need parameter tuning after selecting the suitable model

Some Defination:

a) R-Squared: R-squared is a statistical measure that represents the proportion of the variance for a dependent variable that's explained by an independent variable or variables in a regression model.

If the R2 of a model is 0.50, then approximately half of the observed variation can be explained by the model's inputs.

b) MAE and MSE: MAE : The mean absolute error (MAE) is a measure of errors between paired observations expressing the same phenomenon. MSE : The mean squared error tells you how close a regression line is to a set of points. It does this by taking the distances from the points to the regression line (these distances are the "errors") and squaring them.

c) RMSE: RMSE is a measure of how spread out these residuals are. In other words, it tells you how concentrated the data is around the line of best fit.

In [ ]:

In [ ]:

